

# **REINFORCEMENT LEARNING**

## **PART 3**

Ivan Bratko

Faculty of Computer and Information Science

University of Ljubljana

**ACTIVE LEARNING:**

**TRY TO FIND A GOOD POLICY**

**Ideally: Aim at optimal policy**

# ACTIVE LEARNING

- Active learner attempts to find an optimal policy
- How can ADP be turned into active agent?
  - (1) Agent has to optimise policy, not just execute fixed policy
  - (2) If agent has learned state-transition function  $\delta$  then optimal policy can be determined by solving constraint problem defined by Bellman equations:

$$U(s) = R(s) + \gamma \max_a \sum_{s'} [P(s' | s, a) * U(s')]$$

# ACTIVE AGENT: EXPLOITATION VS. EXPLORATION

- Active agent's task is to search among possible policies
- Actions tried by agent serve two goals:
  - (1) Achieve high reward (**exploitation**)
  - (2) Learning transition model (**exploration**)

# MATHEMATICAL PROBLEM FOR STUDYING OPTIMAL EXPLORATION POLICIES

- What is optimal exploration policy?
- A theoretical framework for studying optimal exploration policies:  
**multi-armed bandit problem**
- n-armed bandit is a gambling machine with n levers (arms - slot machines); gambler inserts a coin, chooses and pulls one of the levers; corresponding machine decides probabilistically about the reward to be paid to the gambler
- Gambler plays many times: what is the best exploration policy (that maximises expected cumulative reward)? Answer in general case is very hard, specially when the machines are not independent
- This is a RL problem

# BALANCE BETWEEN EXPLOITATION AND EXPLORATION

- Is there optimal balance between exploitation and exploration? Main idea is roughly: GLIE
- GLIE schemes try to achieve good balance
- GLIE = **G**reedy in the **L**imit of **E**xploration  
Initially tend to explore, later tend to exploit.  
In the end, method becomes just greedy (choose action that maximizes reward)
- A simple GLIE scheme:  
With probability  $1/t$ , choose random action, otherwise choose highest utility action ( $t$ =time; this probability decreases with time).  
This does converge, but can be very slow.

# TWO MAIN APPROACHES TO ACTIVE RL

Both approaches explore among policies aiming at a good policy, but they differ in what is being learned

1. Utility based learning: Learn optimal state utilities and model delta
2. Q-learning: Learn Q-values

# UTILITY BASED ACTIVE RL

- To search for a good policy, agent carries out trials in the domain
- In choosing next action, the agent strives to good compromise between exploration and exploitation
- The agent updates a domain model, i.e. transition probabilities by counting transitions to next state of form  $(s, a, s')$
- During experimentation, the agent also updates approximations to optimal state utilities and optimal policy
- This can be done by solving the ADP problem; a popular method for that is **value iteration algorithm**



# VALUE ITERATION

- Value iteration is a simple and practical method for determining optimal state utilities, so it is part of common technology for RL
- Given a (current) model delta (i.e. transition probabilities), problem is to solve equations:

$$U(s) = \max_a (P(s' | s,a) * [ r(s,a,s') + \text{gamma} * U(s')]) \quad (\text{Eq. B})$$

- Value iteration method roughly consists of:
  1. Initialise  $U(s)$  for all states  $s$  with arbitrary initial values, e.g.  $U(s)=0$
  2. Keep updating simultaneously all  $U(s)$  according to Eq. B using current  $U(s)$  in right hand side of Eq. B, until difference between old  $U$  and new  $U$  values become sufficiently small

# VALUE ITERATION ALGORITHM

$U'(s)$  are current estimates,  $U(s)$  are estimates from previous iteration

Initialise for all  $s$ :  $U'(s) = 0$ ;

Repeat

for all  $s$ :  $U(s) \leftarrow U'(s)$ ,  $D \leftarrow 0$

for all  $s$  do simultaneously:

$U'(s) \leftarrow \max_a (P(s' | s, a) * [r(s, a, s') + \text{gamma} * U(s')])$

if  $|U'(s) - U(s)| > D$  then  $D \leftarrow |U'(s) - U(s)|$

until  $D < \text{eps} * (1 - \text{gamma}) / \text{gamma}$

**eps** is maximum error allowed in utility of any state

The termination condition comes from mathematical result:

if  $\|U_{i+1} - U_i\| < \text{eps} * (1 - \text{gamma}) / \text{gamma}$  then  $\|U_{i+1} - U\| < \text{eps}$

That is: If successive estimates are close then  $U_{i+1}$  is close to true  $U$

Notation  $\|U\| = \max_s |U(s)|$  (max abs. value in a vector)

# INTUITION WHY SMALL GAMMA HELPS CONVERGENCE

- Consider effects of utilities between distant states:

$S_0 \rightarrow S_1 \rightarrow \dots \rightarrow S_n \rightarrow$


Suppose action in  $S_n$  generates large reward, and actions before  $S_n$  give low rewards. How does this effect  $U(S_0)$  if Gamma is low/high?

# CONVERGENCE OF VALUE ITERATION

- Value iteration algorithm **always converges** to (unique) solution of Bellman equations B.
- In each iteration, error is reduced at least by factor gamma (exponential convergence). However, this gets slow if gamma is close to 1.
- **Convergence to optimal policy is typically faster than convergence to approximation to optimal U values**
- Why is convergence to optimal policy usually faster? Note that current values  $U(s)$  define a policy. When in some iteration  $U(s)$  are sufficiently close to  $U^*(s)$ , these  $U(s)$  define optimal policy  $\pi^*$ . Often  $U(s)$  already define optimal policy even if they are quite far from true values


# VALUE ITERATION, 2x3 GRID EXAMPLE

- A simple robot world consisting of a 2x3 grid. Robot can move horizontally or vertically between adjacent cells; transition into goal state gives reward 100, other transitions reward 0. Gamma = 0.9.



# 2x3 GRID EXAMPLE WITH BROKEN ROBOT

- Let transitions 22-right and 11-right be nondeterministic: 50% chance that robot stays where it is.  $\text{Gamma} = 0.9$ .



	1	2	3
2	0, 45, 62.25, ...	0, 50, 72.5, 82.625 87.18, 89.23, 90.15 90.57, 90.84, 90.88	0
1	0, 0, 0, 81, 81, ...	0, 0, 90, 90	0, 100, 100, ...

# USING OPTIMISTIC UTILITY ESTIMATES

- This is an idea to promote exploration in utility based RL
- Assign higher utilities to states reached by under-explored actions
- Such optimistic utility estimates can be used in value iteration algorithm
- Such optimistic utility estimates promote exploration
- Optimistic estimates can be viewed as reducing desire to explore (curiosity) to (a kind of) greed; high utility estimates reflect underexplored states (novelty of states). A simple greedy policy w.r.t. optimistic estimates will drive the agent to underexplored areas

# OPTIMISTIC ESTIMATES

- $U^+(s)$  = optimistic utility estimate of  $s$  (higher than realistic)
- $U^+(s) \leftarrow \max_a (R(s,a) + \gamma f( \sum_{s'} P(s' | s,a) * U^+(s'), N(s,a) )$
- $N(s,a)$  = # times the pair  $(s,a)$  has occurred so far  
i.e. # times  $a$  has been tried in  $s$
- $f(u, n)$  is called **exploration function**; it trades between greed and curiosity
- To stimulate exploitation,  $f(u,n)$  should increase with  $u$
- To stimulate exploration,  $f(u,n)$  should decrease with  $n$



# AN EXPLORATION FUNCTION

$$f(u,n) = \begin{cases} R^+ & \text{if } n < N_e \\ u & \text{otherwise} \end{cases}$$

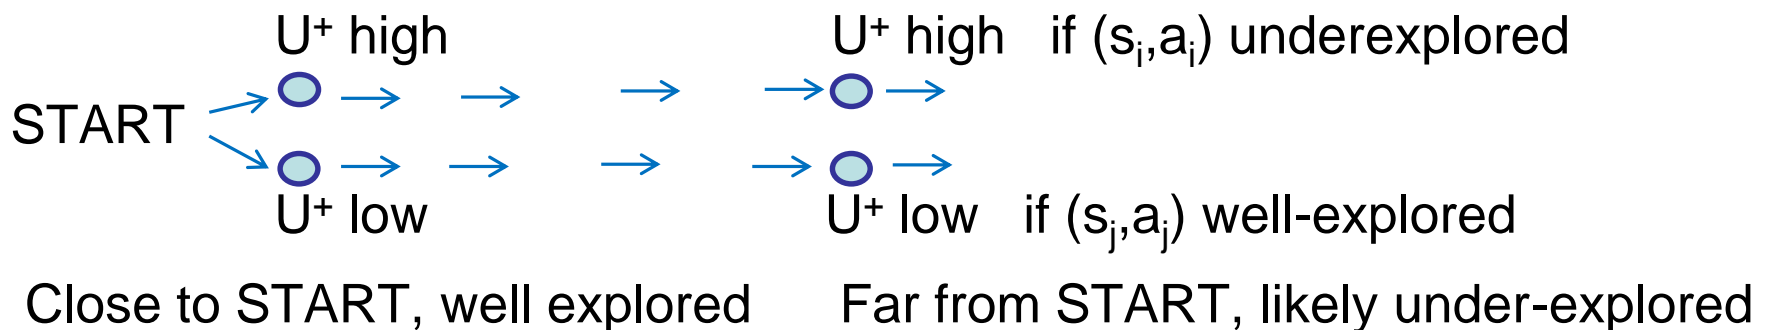
$R^+$  is max. possible cumulative reward obtainable in any state, in practice unrealistically high

$N_e$  is fixed parameter, a threshold for optimism

This  $f$  makes the agent try each action at least  $N_e$  times

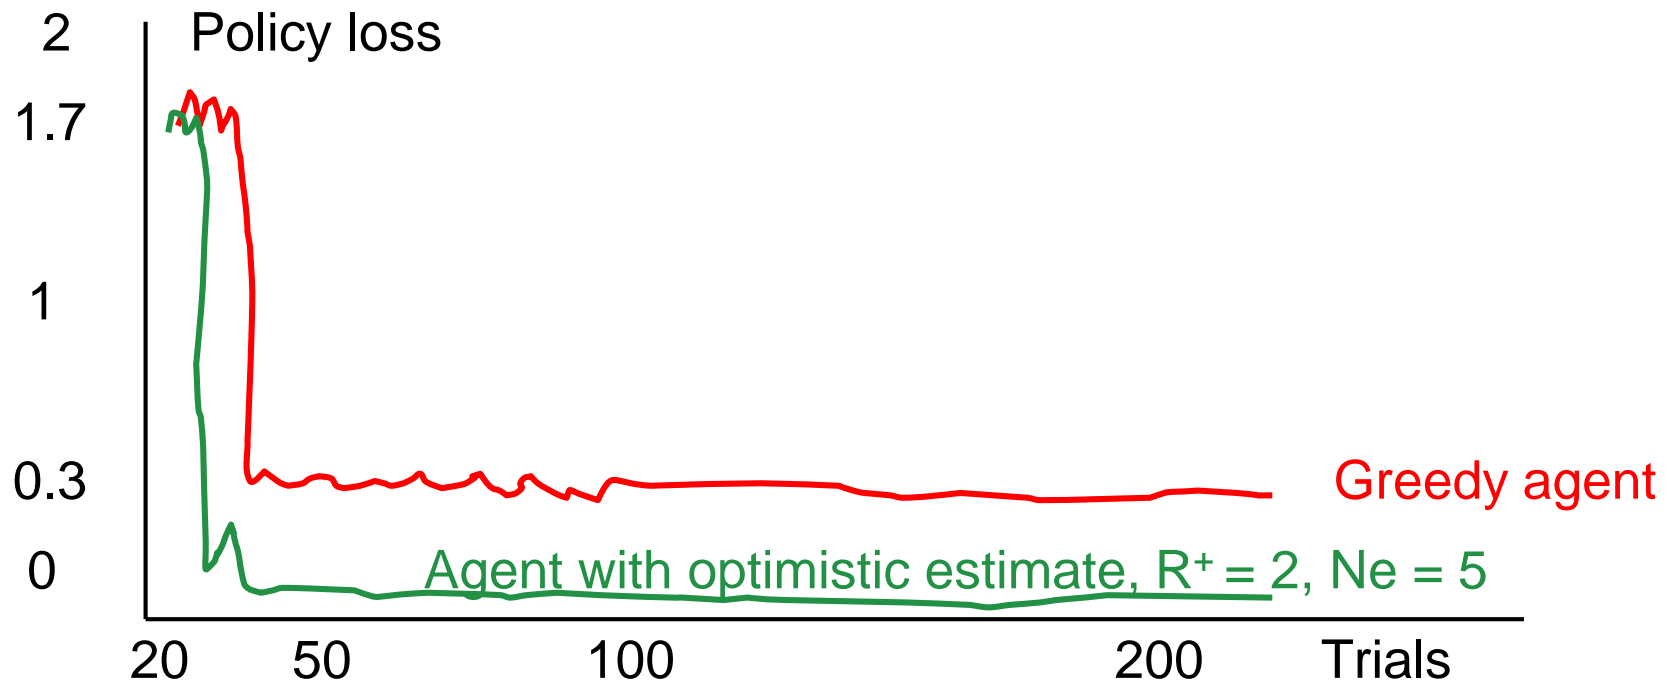
# DEFINITION OF $U^+$

- Note important detail in definition of  $U^+$
- $U^+(s) \leftarrow \max_a (R(s,a) + \gamma \cdot f( \text{SUM } P(\dots) U^+(s'), \dots ))$
- Very important to have  $U^+$  here and not just  $U$
- This drives agent even in states close to START towards remote states in under-explored regions



# RUSSELL & NORVIG 4x3 EXAMPLE

- Convergence of learned policy



Note: Re-drawn approximately from Russell&Norvig

# Q-LEARNING

# Q FUNCTION

- In Q-learning, agent learns function  $Q$ .  $Q(s,a)$  is defined as the maximal cumulative reward achievable by action  $a$  in state  $s$ :

$$Q(s,a) = r(s,a) + \gamma U^*(\delta(s,a)) \quad (\text{for deterministic case})$$

- In utility-based RL, agent has to learn reward function  $r$  and state transition function  $\delta$ . This suffices to determine  $U^*$  function, and that determines the optimal policy: (for deterministic case):

$$PI(s) = \operatorname{argmax}_a [ r(s,a) + \gamma U^*(\delta(s,a)) ]$$

# LEARNING Q FUNCTION

- A TD agent that learns Q function does not need a model of delta
- Therefore: Q-learning is said to be a model-free method
- Complexity comparison of value-based RL and Q-learning: in both cases the domain of functions to be learned are of order  $|S| \times |A|$

# LEARNING Q FUNCTION

- Using identity  $U^*(s) = \max_a Q(s,a)$  gives recursive definition of Q:

$$Q(s,a) = r(s,a) + \gamma \max_{a'} Q(s',a')$$

- One way to learn Q function is through iterative approximations
- $Q^{\wedge}(s,a)$  is current approximation of  $Q(s,a)$
- After each action  $a$  in state  $s$ , approximation is updated by the TD rule:  
$$Q^{\wedge}(s,a) \leftarrow Q^{\wedge}(s,a) + \alpha [ r(s,a) + \gamma \max_{a'} Q^{\wedge}(s',a') - Q^{\wedge}(s,a) ]$$
  
where  $s'$  is result of executing  $a$  in  $s$   
 $\alpha$  = update rate: higher  $\alpha$ , more vigorous update;  $\alpha < 1$

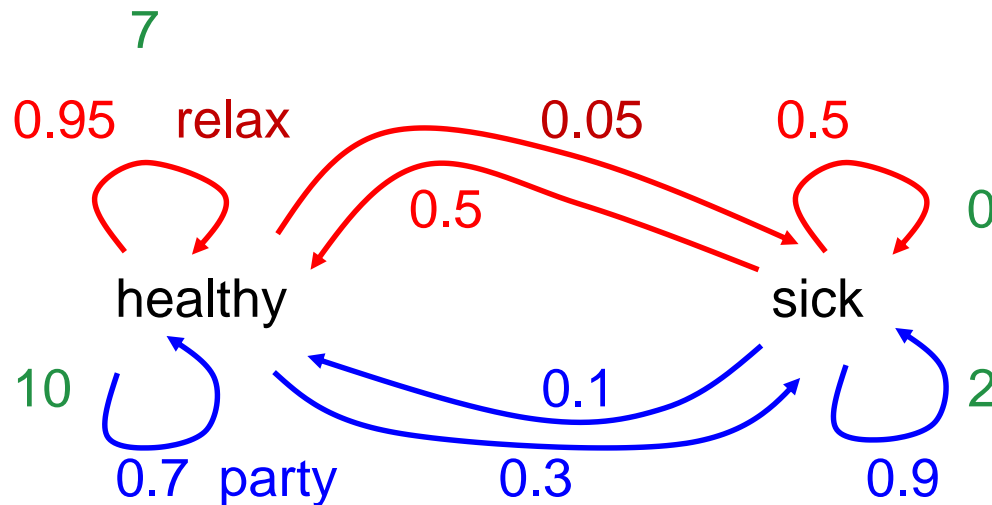
# Q LEARNING ALGORITHM

- For each  $s$  and  $a$ , initialise  $Q^*(s,a) \leftarrow 0$
- Let  $s$  be currently observed state
- Do “forever”:
  - Select an action  $a$  and execute it
  - Receive immediate reward  $r$
  - Observe next state  $s'$
  - Update table entry  $Q^*(s,a)$ :  
 $Q^*(s,a) \leftarrow \dots$  (TD update rule for  $Q^*(s,a)$ )
  - $s \leftarrow s'$



# EXAMPLE FROM POOLE & MACKWORTH 2017

- Sam makes informed decision between partying and relaxing



Green numbers are rewards from actions performed in states

$\text{Gamma} = 0.8$

# WHAT ARE Q-VALUES AFTER THE TRIAL BELOW?

s   a   r   s'	(state,action)	Q(state,action)
-----		
he re 7 he	he, re	$Q^{\wedge}(\text{he, re}) = 0.7 \cdot 0 + 0.3 \cdot (7+1) = 2.1$
he re 7 he	he, re	$Q^{\wedge}(\text{he, re}) = 0.7 \cdot 2.1 + 0.3 \cdot (7+0.8 \cdot 2.1) = 4.07$
he pa 10 he	he, pa	$Q^{\wedge}(\text{he, pa}) = 3.98$
he pa 10 si	he, pa	$Q^{\wedge}(\text{he, pa}) = 5.79$
si pa 2 si	si, pa	...
si re 0 si	si, re	
si, re, 0, he	si, re	

gamma = 0.8, alpha = 0.3

# BEST POLICY FOR SAM

?- utilities(U), q(S,A,Q,U).

U = [healthy/35.714, sick/23.80952],

S = healthy, A = relax, Q = 35.095 ;

S = healthy, A = party, Q = 35.714 ;

S = sick, A = relax, Q = 23.80952 ;

S = sick, A = party, Q = 22.000 ;

# SAM'S Q-LEARNING

After 100 steps (in this domain a trial never ends):

healthy: [relax:34.0675,party:35.4022]

sick: [relax:26.2121,party:22.02939]

After 1000 steps:

healthy: [relax:36.04827,party:35.31005]

sick: [relax:26.5521,party:21.0226]

After 5.000 steps:

healthy: [relax:34.9970,party:36.1772]

sick: [relax:23.5344,party:21.4785]

After 10.000 steps:

healthy: [relax:35.1979,party:36.0068]

sick: [relax:23.7620,party:21.7365]

# RUSSELL & NORVIG ROBOT 4x3

Optimal policy:

In s11: Best action is up,  $Q(s11,up) = 0.7453$

In s31: Best action is left:  $Q(s31,left) = 0.6514$

$Q(s31,up) = 0.6325$

Q-learning, a trial ends in terminal state, or after 100 steps

#trials	$Q(s11,up)$	Best policy
50	0.7989	no
100	0.7788	yes
500	0.7582	yes
1000	0.7716	yes
10000	0.7428	yes

# EXPLORATION STRATEGIES

- Which action to select next?
- Trade-off between exploitation and exploration
- Extreme exploitation: maximise  $_a Q^{\wedge}(s,a)$  (greedy policy)
- This may never discover policies that are even more profitable than the policy already known (also, cf. convergence theorem that requires visiting all  $(s,a)$  pairs many times)
- Epsilon-greedy policy: Choose a random action with probability  $\epsilon$ , and action with maximum utility (according to current estimates) with probability  $1 - \epsilon$ . Parameter  $\epsilon$  may appropriately decrease with time.

# MIXTURE OF EXPLOITATION/EXPLORATION, SOFTMAX

- Make probabilistic choices, higher  $Q^*$  for action  $a_i$ , higher the chance for  $a_i$  to be selected (“softmax”)
- Let  $P(a_i | s)$  be prob. of randomly selecting action  $a_i$  in state  $s$ . Then e.g.

$$P(a_i | s) = \frac{k^{Q^*(s,a_i)}}{\sum_j k^{Q^*(s,a_j)}}$$

$$k > 0$$


higher  $k$  ---> tend to exploit

lower  $k$  ----> tend to explore ( $k \geq 1$ )

E.g.  $Q(s,a_1)=4$ ,  $Q(s,a_2)=1$

$$k=1: P(a_1)=1/(1+1)=0.5, P(a_2)=0.5$$

$$k=2: P(a_1)=2^4/(2^4 + 2^1) = 16/18, P(a_2)=2/18$$



- It is appropriate that  $k$  varies with time: start with desire to explore (small  $k$ ), later prefer to exploit (increase  $k$ )

# CONVERGENCE OF Q LEARNING

- Convergence theorem for Q learning for deterministic MDPs
  - Let all rewards be bounded: for all  $(s,a)$ ,  $|r(s,a)| \leq c$
  - Let agent initialise  $Q^\wedge$  table to arbitrary finite values
  - Let agent use Q learning update rule and discount factor  $\gamma$  s.t.  $0 < \gamma \leq 1$
  - Let for all  $a$  and  $s$ ,  $Q^\wedge_n(s,a)$  denote values of  $Q^\wedge$  after  $n$ -th update
  - Then, if each pair  $(s,a)$  is visited infinitely many times, then  $Q^\wedge_n(s,a)$  converges to  $Q(s,a)$  as  $n \rightarrow \infty$ , for all  $(s,a)$ .



# GENERALISATION OF Q ESTIMATES

- In our methods so far, Q estimates are updated for the visited state-action pairs; these are kept in tabular form, and unvisited values stay unchanged
- This is very limiting
- Therefore other methods determine Q estimates for unvisited state-action pairs through function approximation using ML methods. The estimates for the visited pairs are taken as examples for such ML methods
- Traditionally, neural networks are usually used in this context for function approximation (DRL stands for Deep Reinf. Learning)

# CONVERGENCE OF Q LEARNING IN NONDETERMINISTIC CASE

- Convergence theorem for non-deterministic Q learning
  - Let all rewards be bounded: for all  $(s,a)$ ,  $|r(s,a)| \leq c$
  - Let the agent initialise  $Q^{\wedge}$  values to arbitrary values, and use the **training rule for nondeterministic MDPs** with discount factor  $\gamma$  s.t.  $0 \leq \gamma < 1$
  - Let  $n(i,s,a)$  be the iteration in which action  $a$  was for the  $i$ -th time applied to state  $s$ .
  - If each state-action pair is visited infinitely often,  $0 \leq \alpha_n < 1$ , and  $\sum_{i=1..inf} \alpha_{n(i,s,a)} = inf$  and  $\sum_{i=1..inf} \alpha_{n(i,s,a)}^2 < inf$ , then all  $Q^{\wedge}$  values converge to  $Q$  as  $n \rightarrow inf$ , with probability 1.